
Predicting user ratings with incremental SVD

Stephen Merity

J. Benjamin Cook

Abstract

Hello, World. This is an abstract.

1 Introduction

In this section, we give an introduction to our incremental SVD paper.

2 Dataset

The core of the Netflix dataset consists of 17,770 text files. Each text file represents a distinct movie. The first line in the text file is the movie's unique ID number, which is an integer from 1 to 17,770. All other lines have three comma-delimited entries: user ID, rating, and date.

There are 480,189 unique users in the dataset, however, IDs range from 1 to 2,649,429, with gaps. Ratings are integers from one to five indicating the number of stars the user gave to the movie in question. Dates are in the YYYY-MM-DD format, although we do not use this information in the current project. For example:

```
1 :
1488844,3,2005-09-06
822109,5,2005-05-13
885013,4,2005-10-19
30878,4,2005-12-26
...
```

The above data come from the first movie's text file, which contains ratings for the movie *Dinosaur Planet*. The data indicate that user 1488844 gave *Dinosaur Planet* a three star rating on September 6, 2005, user 822109 gave the movie a five star rating on May 13, 2005, and so on.

In order to be able to perform SVD, we need a matrix with users on the rows and movies on the columns. This matrix would be $480,179 \times 17,770 = 8.5$ billion entries. In a regular matrix format, this would be too big to hold in memory. One estimate is that it takes roughly 65 Gb of RAM to hold the entire matrix (?) although the actual size would depend on the size of the amount of space allocated for each rating. Fortunately, the matrix is extremely sparse, containing only around 100 million non-zero entries. The data structure we use is Python's `scipy.sparse.lil_matrix`. We store data from the text files in this sparse matrix as we read them. After reading in all of the text files, we output the matrix to a Matrix Market format. The Matrix Market format starts with a line containing the dimensions of the matrix and the number of non-zero entries. Then, each line contains $i \ j \ < \text{value} >$. For example, these are the first few lines of a Matrix Market file with a subset of the Netflix data :

```
20000 1000 564726
1 1 3
1 8 4
1 17 2
```

1 30 3
...

Finally, because the process of implementing our incremental SVD system was iterative and because even the iterative method requires serious computational power, we reduced our dataset to smaller subsets for testing. We ran our algorithm on datasets of size 1000×2000 , 1000×3000 ,

3 Method

In the context of recommender systems, the main task of Singular Value Decomposition (SVD) is to decrease the dimensionality of the dataset. We need to be able to summarize the key characteristics of a movie in a much smaller number of features, or variables, than the total number of users in the system. Similarly, we need to reduce the key preferences of users to something much smaller than the total number of movies in the database. The reasons for decreasing the dimensionality of the dataset are two-fold. First, without dimensionality reduction, machine learning tasks would be computationally intractable and second, reducing the dimensionality actually allows us to predict ratings more effectively, since our dataset is so sparse.

3.1 SVD

As with other matrix factorization techniques, SVD works by decomposing a matrix, \mathbf{A} into separate matrices whose product is \mathbf{A} :

$$\text{SVD}(\mathbf{A}) = \mathbf{U}\mathbf{S}\mathbf{V}^T$$

where \mathbf{U} is blank, \mathbf{S} is blank, and \mathbf{V} is blank.

Although we did try to implement SVD from scratch following the algorithm laid out in ?, it quickly became clear that implementing SVD from scratch was beyond the scope of this project.

3.2 Predicting Ratings

3.3 Folding-in

4 Results

These are our results. They are pretty darn good.

5 Discussion

This is a discussion of how well we did.

6 Conclusion

In conclusion, this paper was awesome. Cheers.